

1 | Understanding Authorization

To be able to model and create intuitive and understandable authorization flows, one must understand the basis of authorization flows and the elements by which they are made successful.

In the most simply basis, authorization is the process by which permissions are assigned to a user. There are a few models by which authorization is done, and we will aim to list a few successful systems and their downfalls.

To begin this discussion, we will aim to describe a few terms:

- `model`: a system of authentication, like the ones discussed below
- `agent`: the software tool in a `model` by which authorization is checked
- `rule`: a statement made available to the `agent` to validate claims made by users
- `resource`: a file/page/tool by which the `model` aims to protect
- `action`: what the `agent` grants to do to a `resource`
- `user`: an actor leveraging the `model`'s `agent` to gain perform actions

1.1 | UNIX/BSD PAM

The PAM authentications model, manifested in the `/etc/shadows` files on most `*nix` systems, is one of the most familiar system of authentication to most.

PAM protects individual resources by checking for an octal permission representing whether or not an `action` on a `resource` is accessible to a user or a group of users.

There exists 3 actions: read, write, and execute. The authorization rules are determined on a `resource` level, and `agents` check against rules on access time by users. User permissions override group permissions, which override global permissions.

This systems does not have permission dependencies nor exceptions, group grants and are the only batch executor available.

1.2 | AFS

Access control lists first act separately upon directory resources and file resources. Hence, there are two groups of actions available: those that apply to directories and those that apply to files.

Apart from these permissions, groups/batch executors are applied via lookup and literal expressions-based batched execution by the `agent`. rules, like PAM, are singletons that only grant or remove specific actions.

1.3 | Amazon IAM

Because IAM is the backbone of the Amazon AWS infrastructure, every single lockable AWS permission is considered an `action`. actions could be grouped into presents named "roles", and rules could either apply a role (batched actions) or a specific action to a group.

Inheritance, therefore, could be applied via the creation of groups, applying those permissions to groups, and editing the groups; to prevent conflicts

1.4 | Microsoft Graph Permissions

Microsoft Graph Permissions grants grouped `actions` to specific users. Each `action` have any number (or no) dependencies; hence, when the `agent` applies any `action` in the group to a users, all of the dependents `action` is simultaneously applied.

Similar to PAM, permission groups could be binned and `rules` applied together. However, there are still no user-level dependencies.

1.5 | OAuth

The OAuth `model` allows for authorization of third party `resources` without the `users` credentials.

OAuth creates access tokens which are then passed to third party `models` which use the access token to access `resources` hosted by the resource server.

These temporary access tokens then expire, allowing for a third party to be authenticated without ever giving it credentials and vice versa.

Because of the generalization and lack of `model` of OAuth, there is no hard restriction to permission inheritance

2 | New System Proposal

2.1 | Compiled Features

- PAM is simple!
 - few `actions`
 - groups -> arbitrary user level granularity
 - AFS
 - * Top down from `resources`
 - * resource dependent `actions` (we dont have this!)
 - lack of redundant `actions` (but, we do have this!)
- IAM
 - roles
 - * `actions` are abstracted from entities they apply to
- Microsoft Graphs
 - `actions` can have dependencies on other `actions`
 - * simplicity -> ease of use
 - OAuth
 - * `actions` are arbitrary
 - * authorization is arbitrary
 - auth is abstracted, thus we can protect credentials

2.2 | New System Proposal

Here we propose a new system for authorization that takes some of the most important features from each of the previous standards as highlighted above to create a novel system of authorization that — while increasing slightly in complexity — compensates in terms of its flexibility, elegance, and features.

The core of the system extends the thematic construction of MS Graphs and AWS IAM. Instead of assigning actions or groups of actions, each rule is responsible for tagging a user with a "Role" or a group of Roles.

Each Role is a permissions table containing an arbitrary-length, ordered combination three directives: Add, Subtract, and Link. The Add directive grants a single action to users tagged with that role, the Subtract directives removes a single action to users tagged, and the Link directive includes any other Role's permissions in the hierarchy location of the Link directive.

To Link directive cannot link to any parent nodes, essentially making the chain of permissions of Roles an undirected, fully-connected Tree graph.

When agents parse a user's rules, it reads its list of granted roles in sequential order and performed a DFS of the location to which the role point to on the permissions graph. The results of the DFS would therefore creates a series of new rules which will then be applied sequentially.

The "root" user would therefore be granted the topmost role of the tree by default. Any subordinate users created by root would be manifest in roles which are subsets of the permissions of root by being granted roles that Link to the root role.

This tree-based permissions model essentially implements both the competitive advantage of IAM and Graph. IAM's system of Roles are replicated entirely in this proposed system, except it has the added complexity of trees. Graph's action dependency can be replicated here easily: if two actions are dependent, simply create a Role that encapsulates a Grant for each action, and treating this Role (or inheriting it in other Actions) as the double-dependent action.

Similar to OAuth, all actions except the "Edit All Actions" ("root") action are arbitrarily defined. The root, and any other users who becomes delegated the power, can define actions that then externally becomes validated by the Authentication system (part of the agent) to lock resource.

We have elected to not implement AFS' nature of resource dependent actions. If an action could not be performed upon a resource, the locking resource/status would simply raise a Fault if the user attempts to perform it: hence creating a hidden Authorization scheme.

This scheme are unambiguous as all Grants and Links are fully ordered. An Action cannot be assigned twice in different locations as the permissions Graph is undirected. Groups, similar to the case of PAM, are created by a representative Role that Links to multiple behavioral Roles for the Group and assigning group members that Role.